

GADTs Meet Their Match:

Pattern-Matching Warnings That Account for GADTs, Guards, and Laziness

SEBASTIAN GRAF, Karlsruhe Institute of Technology, Germany

SIMON PEYTON JONES, Microsoft Research, UK

ACM Reference Format:

Sebastian Graf and Simon Peyton Jones. 2020. GADTs Meet Their Match:: Pattern-Matching Warnings That Account for GADTs, Guards, and Laziness. *Proc. ACM Program. Lang.* 1, ICFP, Article 1 (January 2020), 7 pages.

Authors' addresses: Sebastian Graf, Karlsruhe Institute of Technology, Karlsruhe, Germany, sebastian.graf@kit.edu; Simon Peyton Jones, Microsoft Research, Cambridge, UK, simonpj@microsoft.com.

2020. 2475-1421/2020/1-ART1 \$15.00

<https://doi.org/>

Pattern Syntax

$K \in$	Con
$x, y, a, b \in$	Var
$\tau, \sigma \in$	Type
$e \in$	Expr
	$::= x : \tau$
	$ K \bar{a} \bar{y} \bar{e} : \bar{\tau}$
	$ \dots$
$\gamma \in$	TyCt
	$::= \tau_1 \sim \tau_2 \mid \dots$
$g \in$	Grd
	$::= \text{let } x : \tau = e;$
	$ K \bar{a} \bar{y} \bar{y} : \bar{\tau} \leftarrow x;$
	$!x;$

Oracle Syntax

$\Gamma ::= \emptyset \mid \Gamma, x : \tau \mid \Gamma, a$	Context
$\Delta ::= \times \mid \checkmark \mid \delta, \Delta \mid \Delta_1 \vee \Delta_2$	Delta
$\delta ::= \gamma \mid x_1 \approx x_2 \mid K \bar{a} \bar{x} : \bar{\tau} \leftarrow y \mid x \neq K \mid x \approx \perp \mid x \neq \perp \mid \text{let } x = e$	Constraints

Pattern-match Result

$$r ::= \langle \Theta_u, \Theta_d, \Theta_c \rangle$$

$$\langle \Theta_u, \Theta_d, \Theta_c \rangle \cup_U \Theta = \langle \Theta_u \vee \Theta, \Theta_d, \Theta_c \rangle$$

$$\langle \Theta_u, \Theta_d, \Theta_c \rangle \cup_D \Theta = \langle \Theta_u, \Theta_d \vee \Theta, \Theta_c \rangle$$

Constraint tree translation

$$\boxed{\text{ctt } \overline{\text{Grd}} = r}$$

$$\text{ctt } \epsilon = \langle \times, \times, \checkmark \rangle$$

$$\text{ctt } (\text{let } x : \tau = e; \bar{g}) = \text{let } x = e, \text{ctt } \bar{g}$$

$$\text{ctt } (!x; \bar{g}) = (x \neq \perp, \text{ctt } \bar{g}) \cup_D x \approx \perp$$

$$\text{ctt } (K \bar{a} \bar{y} \bar{x} : \bar{\tau} \leftarrow y; \bar{g}) = (K \bar{a} \bar{x} : \bar{\tau} \leftarrow y, \bar{y}, \text{ctt } \bar{g}) \cup_D x \approx \perp \cup_U x \neq K$$

Pattern-match checking

$$\boxed{\text{pmc } \overline{\Delta} \overline{\text{Grd}} = r}$$

$$\text{pmc } \Gamma \epsilon = \langle \times, \times, \checkmark \rangle$$

$$\text{pmc } \Gamma \Delta (\text{let } x : \tau = e; \bar{g}) = \text{pmc } (\Gamma, x : \tau) (\Delta, \text{let } x = e) \bar{g}$$

$$\text{pmc } \Gamma \Delta (!x; \bar{g}) = \text{pmc } \Gamma (\Delta, x \neq \perp) \bar{g}$$

$$\cup_D (\Delta, x \approx \perp)$$

$$\text{pmc } \Gamma \Delta (K \bar{a} \bar{y} \bar{x} : \bar{\tau} \leftarrow y; \bar{g}) = \text{pmc } (\Gamma, \bar{a}, \bar{x} : \bar{\tau}) (\Delta, \bar{y}, K \bar{a} \bar{x} : \bar{\tau} \leftarrow y) \bar{g}$$

$$\cup_D \Delta \oplus x \approx \perp$$

$$\cup_U \Delta \oplus x \neq K$$

1 PROBLEMS WITH CTT

ctt is rather simple now, but it assumes that the incoming Δ is basically unconstrained (e.g. \checkmark). But that certainly is not true for any clause after the first! Intuitively, we replace all leafs in the incoming Δ (which are \checkmark , since we can immediately prune \times). Example:

```


data T = A | B | C
f A A = ()
f B B = ()
f C C = ()

```


We start with $\{(x, y) \mid \Delta\}$ for the uncovered set. After the first clause, we have $A, \checkmark \vee (A \leftarrow x, y \neq A, \checkmark)$ for the uncovered set flowing into the second clause.

The result of `ctt` applied to the second clause is $(x \neq B, \checkmark) \vee (B \leftarrow x, y \neq B, \checkmark)$. But that doesn't consider the incoming uncovered set! For that, we have to substitute every \checkmark in the incoming uncovered set by the constraint tree we just computed.


In tree form. Incoming Δ :



Δ from `ctt` on the second clause:



Substituted into the first Δ :



Note that we now have 4 \checkmark s. So this substitution step reintroduces the exponential blowup.

That alone wouldn't be a problem: Currently, we also would have 4 Δ s in flight for this program. But if we execute on the plan here to separately translate `Grd` into `Con` trees for each clause, and then only *afterwards* (after the substitution step, that is) check for inhabitants (which is conceptually very beautiful), we run into efficiency problems in the implementation, because we have no way to share the work involved with checking the *very similar* branches we just substituted.

It's a lot like choosing the most efficient evaluation strategy, really! Doing the substitution before we digest the tree into a more computably tractable form (like in the current implementation where we cache residual COMPLETE sets) is a lot like call-by-name and we get asymptotically behavior in supposedly trivial cases. The current implementation is more like call-by-value in that regard.


Example, inspired by the test case `ManyAlternatives`:

```

data T = T1 | ... | T1000
f T1 = ()
...
f T1000 = ()


```

The constraint tree of the *covered* set of the 1000th clause will look like this:



The other covered sets are similar. In order to determine whether a clause is redundant, we have to check each of these covered sets for inhabitants! But with COMPLETE sets, we have to constantly check whether the negative constraints form a COMPLETE set. That's very inefficient! And it's the reason we currently have the `vi_cache` field in `VarInfo`: For gradually deleting candidates from the residual COMPLETE sets when we move from clause to clause instead of always beginning from scratch (the full COMPLETE set) at each clause and thinning it out with linearly many negative constraints.

So we definitely want the same kind of caching in our new constraint tree representation. Now here's the problem: I don't currently see how! Intuitively, sharing of work is only possible along the shared path from the root of the final constraint tree we check for inhabitants to one of its leafs. Note how that's not possible in the tree above, because each tree will have a different root, so no sharing on any such paths! If we had the following constraint trees instead:



I.e. with the order of inner nodes reversed, we could share residual COMPLETE sets along the shared path prefix. E.g. the clause for `T500` could re-use the residual COMPLETE sets from `T499`, like it's currently the case.

To achieve this, we have to roll back to the old constraint tree generation scheme, where we pass the incoming Δ to `ctt`.

Test if Oracle state Delta is unsatisfiable

$$\frac{\boxed{\not\vdash_{\text{SAT}} \Gamma \vdash \Delta}}{\not\vdash_{\text{SAT}} \Gamma \vdash fvs\Gamma \triangleright \Delta} \quad \not\vdash_{\text{SAT}} \Gamma \vdash \Delta$$

Test a list of SAT roots for inhabitants

$$\frac{\boxed{\not\vdash_{\text{SAT}} \Gamma \vdash \bar{x} \triangleright \Delta} \quad \not\vdash_{\text{SAT}} \Gamma \vdash x_i \triangleright \Delta}{\not\vdash_{\text{SAT}} \Gamma \vdash \bar{x} \triangleright \Delta}$$

Test a single SAT root for inhabitants

$$\frac{\boxed{\not\vdash_{\text{SAT}} \Gamma \vdash x \triangleright \Delta} \quad \not\vdash_{\text{SAT}} \Gamma \vdash \oplus \Delta x \approx \perp \quad \{\bar{K}\} \text{ COMPLETE set} \quad \overline{\forall \bar{y} : \bar{\tau}. \not\vdash_{\text{SAT}} \Gamma, \bar{y} : \bar{\tau} \vdash \oplus \Delta x \approx K \bar{y}}}{\not\vdash_{\text{SAT}} \Gamma \vdash x \triangleright \Delta}$$

Add a single equality to Delta

$$\boxed{\not\vdash_{\text{SAT}} \Gamma \vdash \oplus \Delta \delta}$$

Term stuff: Bottom, negative info, positive info + generativity, positive info + univalence

$$\frac{x \not\approx sth \in \Delta}{\not\vdash_{\text{SAT}} \Gamma \vdash \oplus \Delta x \approx \perp} \quad \frac{x \approx K \bar{y} \in \Delta}{\not\vdash_{\text{SAT}} \Gamma \vdash \oplus \Delta x \approx \perp}$$

$$\frac{x \not\approx K \in \Delta}{\not\vdash_{\text{SAT}} \Gamma \vdash \oplus \Delta x \approx K \bar{y}} \quad \frac{x \approx K_i \bar{y} \in \Delta \quad i \neq j \quad K_i \text{ and } K_j \text{ generative}}{\not\vdash_{\text{SAT}} \Gamma \vdash \oplus \Delta x \approx K_j \bar{z}}$$

$$\frac{x \approx K \bar{\tau} \bar{y} \in \Delta \quad \not\vdash_{\text{SAT}} \Gamma \vdash \oplus \Delta \tau_i \sim \sigma_i}{\not\vdash_{\text{SAT}} \Gamma \vdash \oplus \Delta x \approx K \bar{\sigma} \bar{z}} \quad \frac{x \approx K \bar{\tau} \bar{y} \in \Delta \quad \not\vdash_{\text{SAT}} \Gamma \vdash \oplus \Delta y_i \approx z_i}{\not\vdash_{\text{SAT}} \Gamma \vdash \oplus \Delta x \approx K \bar{\sigma} \bar{z}}$$

Type stuff: Hand over to unspecified type oracle

$$\frac{\tau_1 \text{ and } \tau_2 \text{ incompatible to Givens in } \Delta \text{ according to type oracle}}{\not\vdash_{\text{SAT}} \Gamma \vdash \oplus \Delta \tau_1 \sim \tau_2}$$

Mixed: Instantiate K and see if that leads to a contradiction TODO: Proper instantiation

$$\frac{\boxed{\not\vdash_{\text{SAT}} \Gamma \vdash y \triangleright \Delta \cup y \not\approx \perp}}{\not\vdash_{\text{SAT}} \Gamma \vdash \oplus \Delta x \approx K \bar{y}}$$

TODO LIST

REFERENCES

- Lennart Augustsson. 1985. Compiling pattern matching. In *Proceedings of the 1985 Conference on Functional Programming and Computer Architecture*.
- Edwin Brady. 2013a. Idris, a general-purpose dependently typed programming language: Design and implementation. *Journal of Functional Programming* 23 (9 2013), 552–593. Issue 05. <https://doi.org/10.1017/S095679681300018X>
- Edwin Brady. 2013b. Programming and Reasoning with Algebraic Effects and Dependent Types. In *Proceedings of the 18th ACM SIGPLAN International Conference on Functional Programming (ICFP '13)*. ACM, New York, NY, USA, 133–144. <https://doi.org/10.1145/2500365.2500581>
- James Cheney and Ralf Hinze. 2003. *First-class phantom types*. Technical Report. Cornell University.
- Koen Claessen, Moa Johansson, Dan Rosén, and Nicholas Smallbone. 2013. Automating Inductive Proofs Using Theory Exploration.. In *CADE (Lecture Notes in Computer Science)*, Maria Paola Bonacina (Ed.), Vol. 7898. Springer, 392–406.
- Thierry Coquand. 1992. Pattern matching with dependent types. In *Proceedings of the Workshop on Types for Proofs and Programs*.
- Joshua Dunfield. 2007a. Refined Typechecking with Stardust. In *Proceedings of the 2007 Workshop on Programming Languages Meets Program Verification (PLPV '07)*. ACM, New York, NY, USA, 21–32. <https://doi.org/10.1145/1292597.1292602>
- Joshua Dunfield. 2007b. *A Unified System of Type Refinements*. Ph.D. Dissertation. Carnegie Mellon University. CMU-CS-07-129.
- Richard A. Eisenberg and Stephanie Weirich. 2012. Dependently Typed Programming with Singletons. In *Proceedings of the 2012 Haskell Symposium (Haskell '12)*. ACM, New York, NY, USA, 117–130. <https://doi.org/10.1145/2364506.2364522>
- M Erwig and SL Peyton Jones. 2000. Pattern guards and transformational patterns. In *Proceedings of the 2000 Haskell Symposium*. ACM.
- Jacques Garrigue and Jacques Le Normand. 2011. Adding GADTs to OCaml: the direct approach. In *Workshop on ML*.
- Jean-Yves Girard, Paul Taylor, and Yves Lafont. 1989. *Proofs and Types*. Cambridge University Press, New York, NY, USA.
- Georgios Karachalias, Tom Schrijvers, Dimitrios Vytiniotis, and Simon Peyton Jones. 2015. *GADTs meet their match (extended version)*. Technical Report. KU Leuven. http://people.cs.kuleuven.be/~george.karachalias/papers/gadtpm_ext.pdf
- Neelakantan R. Krishnaswami. 2009. Focusing on Pattern Matching. In *Proceedings of the 36th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL '09)*. ACM, New York, NY, USA, 366–378. <https://doi.org/10.1145/1480881.1480927>
- Alain Laville. 1991. Comparison of Priority Rules in Pattern Matching and Term Rewriting. *J. Symb. Comput.* 11, 4 (May 1991), 321–347. [https://doi.org/10.1016/S0747-7171\(08\)80109-5](https://doi.org/10.1016/S0747-7171(08)80109-5)
- Fabrice Le Fessant and Luc Maranget. 2001. Optimizing Pattern-Matching. In *Proceedings of the 2001 International Conference on Functional Programming*.
- Luc Maranget. 1992. Compiling Lazy Pattern Matching. In *Proceedings of the 1992 ACM Conference on LISP and Functional Programming (LFP '92)*. ACM, New York, NY, USA, 21–31. <https://doi.org/10.1145/141471.141499>
- Luc Maranget. 2007. Warnings for pattern matching. *Journal of Functional Programming* 17 (2007), 387–421. Issue 3.
- Luc Maranget. 2008. Compiling pattern matching to good decision trees. In *Proceedings of the ACM Workshop on ML*.
- Luc Maranget and Projet Para. 1994. *Two Techniques for Compiling Lazy Pattern Matching*. Technical Report.
- The Coq development team. 2004. *The Coq proof assistant reference manual*. LogiCal Project. <http://coq.inria.fr> Version 8.0.
- C. McBride and J. McKinna. 2004. The view from the left. *Journal of Functional Programming* 14, 1 (2004), 69–111.
- Neil Mitchell and Colin Runciman. 2008. Not All Patterns, but Enough: An Automatic Verifier for Partial but Sufficient Pattern Matching. In *Proceedings of the First ACM SIGPLAN Symposium on Haskell (Haskell '08)*. ACM, New York, NY, USA, 49–60. <https://doi.org/10.1145/1411286.1411293>
- Ulf Norell. 2007. *Towards a practical programming language based on dependent type theory*. Ph.D. Dissertation. Department of Computer Science and Engineering, Chalmers University of Technology, SE-412 96 Göteborg, Sweden.
- Ulf Norell. 2008. Dependently typed programming in Agda. In *Lecture Notes from the Summer School in Advanced Functional Programming*.
- Simon Peyton Jones, Dimitrios Vytiniotis, Stephanie Weirich, and Geoffrey Washburn. 2006. Simple Unification-based Type Inference for GADTs. In *Proceedings of the Eleventh ACM SIGPLAN International Conference on Functional Programming (ICFP '06)*. ACM, New York, NY, USA, 50–61. <https://doi.org/10.1145/1159803.1159811>
- Norman Ramsey, João Dias, and Simon Peyton Jones. 2010. Hoopl: A Modular, Reusable Library for Dataflow Analysis and Transformation. In *Proceedings of the Third ACM Haskell Symposium on Haskell (Haskell '10)*. ACM, New York, NY, USA, 121–134. <https://doi.org/10.1145/1863523.1863539>
- Patrick M. Rondon, Ming Kawaguci, and Ranjit Jhala. 2008. Liquid Types. In *Proceedings of the 2008 ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI '08)*. ACM, New York, NY, USA, 159–169. <https://doi.org/10.1145/1375581.1375602>
- Tom Schrijvers, Simon Peyton Jones, Manuel Chakravarty, and Martin Sulzmann. 2008. Type Checking with Open Type Functions. In *Proceedings of the 13th ACM SIGPLAN International Conference on Functional Programming (ICFP '08)*. ACM,

- New York, NY, USA, 51–62. <https://doi.org/10.1145/1411204.1411215>
- Tom Schrijvers, Simon Peyton Jones, Martin Sulzmann, and Dimitrios Vytiniotis. 2009. Complete and Decidable Type Inference for GADTs. In *Proceedings of the 14th ACM SIGPLAN International Conference on Functional Programming (ICFP '09)*. ACM, New York, NY, USA, 341–352. <https://doi.org/10.1145/1596550.1596599>
- R. C. Sekar, R. Ramesh, and I. V. Ramakrishnan. 1995. Adaptive Pattern Matching. *SIAM J. Comput.* 24, 6 (Dec. 1995), 1207–1234. <https://doi.org/10.1137/S0097539793246252>
- Peter Sestoft. 1996. ML pattern match compilation and partial evaluation. In *Partial Evaluation*, Olivier Danvy, Robert Glück, and Peter Thiemann (Eds.). Lecture Notes in Computer Science, Vol. 1110. Springer Berlin Heidelberg, 446–464. https://doi.org/10.1007/3-540-61580-6_22
- Tim Sheard. 2004. Languages of the Future. In *In OOPSLA '04: Companion to the 19th annual ACM SIGPLAN conference on Object-oriented programming systems, languages, and applications*. ACM Press, 116–119.
- W. Sonnex, S. Drossopoulou, and S. Eisenbach. 2012. Zeno: An Automated Prover for Properties of Recursive Data Structures. Springer-Verlag Berlin, 407–421. https://doi.org/10.1007/978-3-642-28756-5_28
- Martin Sulzmann, Manuel M. T. Chakravarty, Simon Peyton Jones, and Kevin Donnelly. 2007. System F with Type Equality Coercions. In *Proceedings of the 2007 ACM SIGPLAN International Workshop on Types in Languages Design and Implementation (TLDI '07)*. ACM, New York, NY, USA, 53–66. <https://doi.org/10.1145/1190315.1190324>
- Peter Thiemann. 1993. Avoiding Repeated Tests in Pattern Matching. In *3rd International Workshop on Static Analysis*, Gilberto Filé (Ed.). Padova, Italia, 141–152.
- Niki Vazou, Eric L. Seidel, Ranjit Jhala, Dimitrios Vytiniotis, and Simon Peyton-Jones. 2014. Refinement Types for Haskell. In *Proceedings of the 19th ACM SIGPLAN International Conference on Functional Programming (ICFP '14)*. ACM, New York, NY, USA, 269–282. <https://doi.org/10.1145/2628136.2628161>
- Dimitrios Vytiniotis, Simon Peyton Jones, Tom Schrijvers, and Martin Sulzmann. 2011. Outsidein(x) Modular Type Inference with Local Assumptions. *J. Funct. Program.* 21, 4-5 (Sept. 2011), 333–412. <https://doi.org/10.1017/S0956796811000098>
- Philip Wadler. 1987a. Efficient compilation of pattern matching. In *The implementation of functional programming languages*, SL Peyton Jones (Ed.). Prentice Hall, 78–103.
- P. Wadler. 1987b. Views: A Way for Pattern Matching to Cohabit with Data Abstraction. In *Proceedings of the 14th ACM SIGACT-SIGPLAN Symposium on Principles of Programming Languages (POPL '87)*. ACM, New York, NY, USA, 307–313. <https://doi.org/10.1145/41625.41653>
- Hongwei Xi. 1998a. Dead Code Elimination Through Dependent Types. In *Proceedings of the First International Workshop on Practical Aspects of Declarative Languages (PADL '99)*. Springer-Verlag, London, UK, 228–242.
- Hongwei Xi. 1998b. *Dependent Types in Practical Programming*. Ph.D. Dissertation. Carnegie Mellon University.
- Hongwei Xi. 2003. Dependently typed pattern matching. *Journal of Universal Computer Science* 9 (2003), 851–872.
- Hongwei Xi, Chiyen Chen, and Gang Chen. 2003. Guarded Recursive Datatype Constructors. In *Proceedings of the 30th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL '03)*. ACM, New York, NY, USA, 224–235. <https://doi.org/10.1145/604131.604150>
- Dana N. Xu. 2006. Extended Static Checking for Haskell. In *Proceedings of the 2006 ACM SIGPLAN Workshop on Haskell (Haskell '06)*. ACM, New York, NY, USA, 48–59. <https://doi.org/10.1145/1159842.1159849>
- Dana N. Xu, Simon Peyton Jones, and Koen Claessen. 2009. Static Contract Checking for Haskell. In *Proceedings of the 36th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL '09)*. ACM, New York, NY, USA, 41–52. <https://doi.org/10.1145/1480881.1480889>
- Brent A. Yorgey, Stephanie Weirich, Julien Cretin, Simon Peyton Jones, Dimitrios Vytiniotis, and José Pedro Magalhães. 2012. Giving Haskell a Promotion. In *Proceedings of the 8th ACM SIGPLAN Workshop on Types in Language Design and Implementation (TLDI '12)*. ACM, New York, NY, USA, 53–66. <https://doi.org/10.1145/2103786.2103795>